


Day - 1

Machine Learning Refresher

What is Machine Learning (ML) ?

→ Machine Learning (ML) is a subset of **Artificial Intelligence (AI)** that allows machines (computers or systems) to **learn** from data without being explicitly (**programmatically**) programmed. In simpler terms, it's about teaching machines how to recognize patterns and make decisions based on data, rather than giving them fixed instructions for every task.

→ The machine uses past experiences (data) to identify patterns, make predictions, and even improve its performance over time without being told exactly how to do so.



Types of Machine Learning ?

1. Supervised Machine Learning

→ In **Supervised Learning**, the model is trained using labeled data, where the input data is paired with the correct output (label). The model learns to map inputs to outputs based on this data. The goal is for the model to make accurate predictions when new, unseen data is introduced.

How it Works:

- The system is provided with input-output pairs (training data).
- The model makes predictions.
- The system is corrected based on the difference between its prediction and the actual result (feedback).
- This process continues until the model's accuracy improves.

Real-World Examples:

- **Email Spam Detection:** The algorithm is trained with labeled examples of spam and non-spam emails. It learns patterns to classify new emails as spam or not.
- **Image Classification:** Given labeled images (e.g., pictures of cats and dogs), the model learns to classify new images as either a cat or dog.
- **Predicting House Prices:** The model is trained on past data (e.g., house features like size, location, number of rooms, etc.) to predict the price of a new house.

Types of Unsupervised Learning Algorithms:

- **Linear Regression:** Used for predicting continuous values (e.g., predicting the price of a house).
 - **Logistic Regression:** Used for binary classification (e.g., predicting whether an email is spam or not).
 - **Decision Trees:** A model that splits the data into branches to make decisions.
 - **Random Forests:** An ensemble method that combines multiple decision trees to improve accuracy.
 - **Support Vector Machines (SVM):** Used for classification and regression tasks.
 - **Neural Networks:** Mimic the human brain to make predictions, especially useful for complex problems like image and speech recognition.
-

2. Unsupervised Machine Learning

→ In **Unsupervised Learning**, the model is trained using data that is not labeled. The system tries to learn patterns or structures from the data itself without prior knowledge of what the output should be.

How it Works:

- The system receives input data without any labels (i.e., there's no "correct" answer given).
- The goal is for the algorithm to discover inherent patterns or relationships in the data, such as grouping similar data points or reducing the dimensions of the data.

Real-World Examples:

- **Customer Segmentation:** In marketing, unsupervised learning is used to segment customers into different groups based on behaviors, preferences, or demographics. These groups can then be targeted with specific marketing strategies.
- **Anomaly Detection:** Used in fraud detection, where the algorithm looks for unusual patterns in data (e.g., unusual spending patterns in credit card transactions).
- **Topic Modeling:** Unsupervised learning can be used to group similar documents or articles based on their topics (e.g., grouping news articles into categories like politics, sports, or entertainment).

Types of Unsupervised Learning Algorithms:

- **Clustering:** Groups similar data points together. Example: **K-means clustering, DBSCAN.**
 - **Dimensionality Reduction:** Reduces the number of features in the data while preserving important information. Example: **Principal Component Analysis (PCA), t-SNE.**
 - **Association Rule Learning:** Finds relationships or patterns in datasets, such as items that are frequently bought together. Example: **Apriori Algorithm.**
-

3. Reinforcement Machine Learning

→ In **Reinforcement Learning**, the model learns by interacting with an environment and receiving feedback based on its actions. Unlike supervised learning, where the model is trained with labeled data, RL systems learn by trial and error, trying to maximize rewards over time.

How it Works:

- An agent (the model) takes actions in an environment.
- For each action, the agent receives a reward or penalty.
- The goal is for the agent to learn the optimal strategy (policy) that maximizes its cumulative reward over time.

Real-World Examples:

- **Game Playing (Chess, Go, or Video Games):** In games, RL is used to teach an agent to play the game by rewarding it when it makes good moves and penalizing it for bad ones. This is how AI programs like AlphaGo and AlphaZero were able to defeat world champions in complex games.
- **Robotics:** Robots use RL to learn how to perform tasks, such as navigating through a room or assembling objects. The robot receives rewards for taking correct actions and penalties for incorrect ones.
- **Self-Driving Cars:** RL is used to teach autonomous vehicles to navigate through traffic, obey traffic laws, and avoid accidents by learning from interactions with the environment.

Types of Reinforcement Learning Algorithms:

- **Q-Learning:** A popular algorithm used for learning the value of actions in a given state.
- **Deep Q-Networks (DQN):** Combines deep learning and Q-learning for more complex tasks.
- **Policy Gradient Methods:** These methods directly learn the policy (a mapping from states to actions).
- **Monte Carlo Tree Search (MCTS):** Used for decision-making in environments like games

What is Predictive Modeling?

→ **Predictive Modeling** refers to the process of creating, testing, and validating a model that is used to predict future outcomes based on historical data. It is a technique used to forecast values or trends based on patterns found in data.

While **predictive modeling** is not inherently tied to deep learning, deep learning can be a powerful tool for predictive modeling, especially when dealing with complex datasets.

Key Concepts in Predictive Modeling:

1. **Input Variables (Features):** These are the characteristics or attributes used to make predictions. For example, in predicting house prices, features could include square footage, location, number of bedrooms, etc.
2. **Output Variable (Target):** This is the value you're trying to predict, like the price of a house, the likelihood of an event, or a future sales number.
3. **Model:** The mathematical representation that takes input variables and predicts the output. The model can be a linear regression model, decision tree, random forest, or even a deep learning model.
4. **Training and Validation:** The model is trained using historical data, and its accuracy is validated by testing it against unseen data.

Types of Predictive Models:

- **Linear Regression:** Predicts continuous values by finding the relationship between input features and output.
- **Logistic Regression:** Used for binary classification problems (e.g., predicting whether a customer will churn or not).
- **Decision Trees and Random Forests:** Used for classification and regression tasks by breaking data into decision nodes based on feature values.
- **Support Vector Machines (SVM):** Used for classification and regression by finding the hyperplane that best separates data points.
- **Neural Networks/Deep Learning:** Used for highly complex data like images, text, or speech, where traditional models may not perform as well.

Real-World Applications of Predictive Modeling:

1. **Customer Churn Prediction:** Predicting which customers are likely to leave a service, allowing businesses to take preemptive action.
2. **Financial Forecasting:** Predicting stock prices, sales trends, or revenue for a company.
3. **Healthcare:** Predicting disease outcomes, patient readmissions, or responses to treatment.
4. **Fraud Detection:** Identifying fraudulent transactions by recognizing patterns in transaction data.
5. **Risk Management:** Predicting risks in areas like insurance, loans, and manufacturing to prevent losses.

What is Deep Learning ?

→ **Deep Learning** is a subset of **Machine Learning (ML)** that deals with algorithms inspired by the human brain's structure, called **Artificial Neural Networks (ANNs)**. These models are designed to automatically learn from data in a way that mimics the way humans learn through experience.

→ Deep learning uses **multi-layered neural networks** to model complex relationships and patterns in large amounts of data. The more layers in the network, the deeper the learning model, which gives it the name "deep learning." Deep learning has revolutionized fields like **computer vision**, **natural language processing**, **speech recognition**, and even **gaming** due to its ability to handle complex data and tasks.

How Does Deep Learning Work?

At the core of deep learning is the concept of a **neural network**. This network consists of layers of nodes (also known as **neurons**), and each node is responsible for learning features or patterns in the data.

Basic Components of a Neural Network:

1. **Input Layer:** This is where the data enters the network. For example, in an image classification task, each pixel of the image will be an input feature.
2. **Hidden Layers:** These layers consist of nodes (neurons) that process the data. There can be many hidden layers in a deep learning model. Each layer processes the input and sends the output to the next layer.
3. **Output Layer:** This layer produces the result or prediction. For example, in an image classification task, the output layer could contain categories like "cat," "dog," or "car."
4. **Weights and Biases:** These are parameters that help the model learn. During the training process, the model adjusts the weights to improve its accuracy.

Types of Deep Learning Networks

1. Convolutional Neural Networks (CNNs)

- **Purpose:** Primarily used for tasks involving images or videos.
- **How it Works:** CNNs apply **convolutional layers** that help the model focus on small, local patterns in the data (e.g., edges in an image). The layers gradually build up to more complex features like shapes and objects.

Real-Life Example:

- **Image Classification:** A CNN can be trained to recognize images of cats and dogs. It will first learn basic patterns like edges, then shapes, and eventually complex features like fur or ears to distinguish between cats and dogs.

2. Recurrent Neural Networks (RNNs)

- **Purpose:** Used for sequential data (data with a time or order component).
- **How it Works:** RNNs use feedback loops to remember information from previous time steps. They are good at understanding sequences, like sentences, speech, or time series.

Real-Life Example:

- **Speech Recognition:** An RNN can be used to convert spoken language into text, as it can understand the context and order of words in a sentence.
- **Stock Price Prediction:** RNNs can be trained on historical stock data to predict future prices based on trends.

3. Generative Adversarial Networks (GANs)

- **Purpose:** Used to generate new data that mimics the distribution of the original data.
- **How it Works:** GANs consist of two parts:**Real-Life Example:**
 - **Generator:** Creates fake data.
 - **Discriminator:** Tries to distinguish between real and fake data.

- **Image Generation:** GANs can generate realistic images from scratch, like creating new faces or generating art that looks like it was painted by a human.

4. Transformers (Self-Attention Mechanisms)

- **Purpose:** Used mainly for tasks related to language (e.g., text generation, translation).
- **How it Works:** Unlike RNNs, Transformers use an **attention mechanism** to focus on different parts of the input sequence. This allows them to handle long-range dependencies better and more efficiently.

Real-Life Example:

- **Machine Translation:** Google Translate uses Transformer-based models to translate text between languages.
- **Text Generation:** GPT (Generative Pre-trained Transformer) models like ChatGPT use transformers to generate coherent and contextually relevant text.

What is Correlation?

→ **Correlation** is a statistical term that describes the relationship or connection between two (or more) variables. It tells us how one variable changes in relation to another. The basic idea is that if one variable increases or decreases, does the other one follow the same trend? Or does it move in the opposite direction? Or is there no predictable relationship between them?

In simple terms, correlation is about understanding how **two things are related** to each other.

Types of Correlation

There are three types of correlation, which describe different kinds of relationships between variables:

1. Positive Correlation:

- When one variable increases, the other also increases.
- For example, there's a positive correlation between **study time** and **exam scores**. The more time you spend studying, the higher your exam scores are likely to be.

Real-Life Example:

- **Height and Weight:** In general, the taller a person is, the more they tend to weigh. So, height and weight have a positive correlation.

2. Negative Correlation:

- When one variable increases, the other decreases.
- For example, there's a negative correlation between **exercise frequency** and **body fat percentage**. The more you exercise, the lower your body fat percentage tends to be.

Real-Life Example:

- **Temperature and Heating Costs:** As the **temperature** goes up (during summer), **heating costs** go down. So, they have a negative correlation.

3. Zero or No Correlation:

- When there is no predictable relationship between the two variables. Changes in one variable do not affect the other.

Real-Life Example:

- **Shoe Size and Intelligence:** There is no predictable relationship between a person's shoe size and their intelligence. Hence, the correlation between shoe size and intelligence is close to zero.

How is Correlation Measured?

Correlation is typically measured using a statistic called **Pearson's correlation coefficient (r)**, which ranges from **-1 to 1**.

- **+1** indicates a **perfect positive correlation**: As one variable increases, the other always increases at a constant rate.

- **1** indicates a **perfect negative correlation**: As one variable increases, the other always decreases at a constant rate.
- **0** indicates **no correlation**: No relationship exists between the two variables.
- **Between 0 and 1** or **0 and -1**: The correlation is somewhere between perfect positive or perfect negative.

Here's how you interpret Pearson's correlation coefficient:

- **0.7 to 1**: Strong positive correlation
 - **0.3 to 0.7**: Moderate positive correlation
 - **0 to 0.3**: Weak positive correlation
 - **0.3 to 0**: Weak negative correlation
 - **0.7 to -1**: Strong negative correlation
 - **0**: No correlation
-

Examples of Correlation in Real Life

Example 1: Sales and Advertising

- **Scenario**: A company spends money on advertising. Over time, they observe that when they increase advertising spending, their sales go up.
- **Correlation**: There's a positive correlation between **advertising spending** and **sales**.
- **Interpretation**: The more money the company spends on advertising, the higher their sales are likely to be.

Example 2: Age and Income

- **Scenario**: People often experience a rise in income as they get older due to career progression.
- **Correlation**: There's generally a **positive correlation** between **age** and **income**, especially for working professionals in established careers.
- **Interpretation**: As people age and gain more experience, their income tends to rise.

Example 3: Temperature and Ice Cream Sales

- **Scenario**: On hotter days, ice cream sales tend to rise.
- **Correlation**: There's a positive correlation between **temperature** and **ice cream sales**.
- **Interpretation**: The hotter the weather, the more ice cream is likely to be sold.

Example 4: Education Level and Job Satisfaction

- **Scenario**: Studies show that people with higher education levels tend to report higher job satisfaction.
- **Correlation**: There may be a **positive correlation** between **education level** and **job satisfaction**.
- **Interpretation**: People with more education might have jobs that are more aligned with their interests or offer better pay, leading to greater job satisfaction.

Example 5: Exercise and Sleep Quality

- **Scenario**: Research shows that regular exercise can improve the quality of sleep.
- **Correlation**: There's a **positive correlation** between **exercise frequency** and **sleep quality**.
- **Interpretation**: The more you exercise, the better your sleep quality tends to be.

Example 6: Education and Income (with a Negative Aspect)

- **Scenario**: In some situations, higher education might correlate with higher income, but not always. Some people might be overqualified for a job, leading to dissatisfaction and lower income than expected.
- **Correlation**: The relationship might be **negative** in certain situations because people with higher education levels might take lower-paying jobs or experience frustration from not using their full skills.
- **Interpretation**: The relationship between education and income could vary based on the field of study and career choices.

How to Visualize Correlation:

To visualize correlation, you can use a **scatter plot**, which is a graph where each point represents a pair of values for two variables.

- **Positive Correlation:** In a scatter plot, the points will generally go from the bottom left to the top right (an upward trend).
- **Negative Correlation:** The points will go from the top left to the bottom right (a downward trend).
- **Zero Correlation:** The points will appear scattered randomly with no clear trend.

Example of Scatter Plots:

- **Positive Correlation:**

If you plot data for **hours studied** vs. **exam score**, the plot would show an upward trend (more study time → higher exam scores).

- **Negative Correlation:**

If you plot **temperature** vs. **heating costs**, you'd see a downward trend (higher temperature → lower heating costs).